

Lab Course Machine Learning Exercise Sheet 1

Python Warmup

Question 1

```
In [ ]: import csv
import codecs
import urllib.request
from collections import Counter
import glob
import codecs
import re
import pandas as pd
import math
from cmath import exp
import matplotlib.pyplot as plt
import numpy as np
import nltk
from nltk.corpus import stopwords
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from pandasql import sqldf
import pandas as pd
from sklearn import datasets
```

[2 points] In this part of the assignment, you have to write a word count program. Your program should read the provided text document on learnweb named random text.txt and then output the following stats:

```
In [ ]: filepath = "/Users/farjad.ahmed/Documents/Studies/ML Lab/Exercise_01/random_text.txt"
with open(filepath) as f:
    sentences= f.readlines()
```

```
In [ ]: # print(stopwords.words('english'))
```

```
In [ ]: bigString = ''
        for sent in sentences:
            bigString = bigString + sent

        import string
        bigString = bigString.translate(str.maketrans('', '', string.punctuation))
        tokens = nltk.word_tokenize(bigString)
```

```
In [ ]: stop_words = set(stopwords.words('english'))
        word_tokens = nltk.word_tokenize(bigString)
        sentence_filtered = [w for w in word_tokens if not w.lower() in stop_words]
        sentence_filtered = []
        for word in word_tokens:
            if word not in stop_words:
                sentence_filtered.append(word)
```

Part a

a) The number of unique non-stop words. (Hint: you can use "nltk" library to get a list of English language stop words.)

```
In [ ]: len(sentence_filtered)
```

```
Out[ ]: 615
```

```
In [ ]: # Finding the unique set of words that do not repeat
        len(list(set(sentence_filtered)))
```

```
Out[ ]: 403
```

Part b

b) The top 5 most frequent non-stop words.

```
In [ ]: from collections import Counter
        myDict = Counter(sentence_filtered)
        countList = list(myDict.items())
        getList = sorted(countList, key=lambda x: x[1], reverse=True)
        getList[:5]
```

```
Out[ ]: [('Harry', 26), ('Voldemort', 9), ('also', 8), ('He', 8), ('Dark', 7)]
```

Question 2

[2 points] In a simple regression problem we fit a straight line $y = b_1x + b_0$ to a given data. However, not all problems in nature are by default linear. Given the data below see if a straight line is a good fit.

In cases where the data does not follow a linear trend, one can transform the variables and then apply the linear regression technique to better fit the data. From the given choices, try which function would be a better representation for the data.

(a) Linear : $y = b_1x + b_0$ (b) Power : $y = bx^m$ (c) Exponential : $y = be^{mx}$ (d) Logarithmic : $y = m\log x + b$ (e) Reciprocal : $y = 1/(mx+b)$

Generate a 2 x 2 subplot with the following techniques, plot, semilogx, semilogy, loglog. Read about these plotting techniques. These plots will let you understand which of the above 5 choices will be the best fit. Plot the data points and the best fit curve in a well-formatted plot with axis labels, title and the legend. (Hint: you can use the polyfit function from numpy for this part.)

Plotting all equations for all plotting methods. Answer: Power Estimates the data the best, it is evident from the working below.

```
In [ ]: d = {'x': [0.0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0], 'y': [6.0, 4.83, 3.7, 3.15, 2.41, 1.83, 1.49, 1.21, 0.96, 0.8, 0.6]}
df = pd.DataFrame(data=d)
```

```
In [ ]: plt.rcParams['figure.figsize'] = [15, 15]
```

```
In [ ]: x_2 = df['x'].to_numpy()
y_1 = df['y'].to_numpy()
xm = np.mean(x_2)
ym = np.mean(y_1)
b1 = np.sum((x_2-xm)*(y_1-ym))/(np.sum((x_2-xm)**2))
b0 = ym - b1*xm
print('b1: {} and b0: {}'.format(b1,b0))
```

```
b1: -1.022181818181818 and b0: 5.005454545454546
```

```
In [ ]: m = b1 #gradient
b = b0 #intercept
```

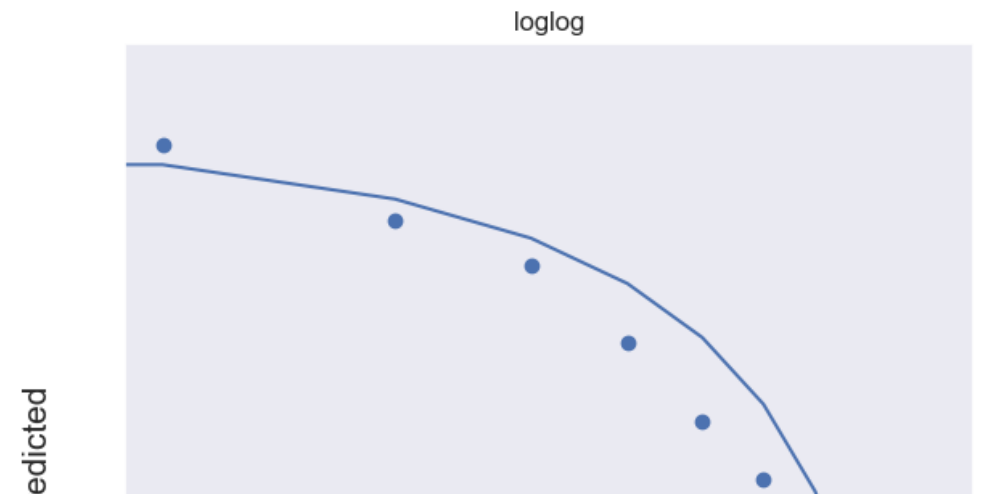
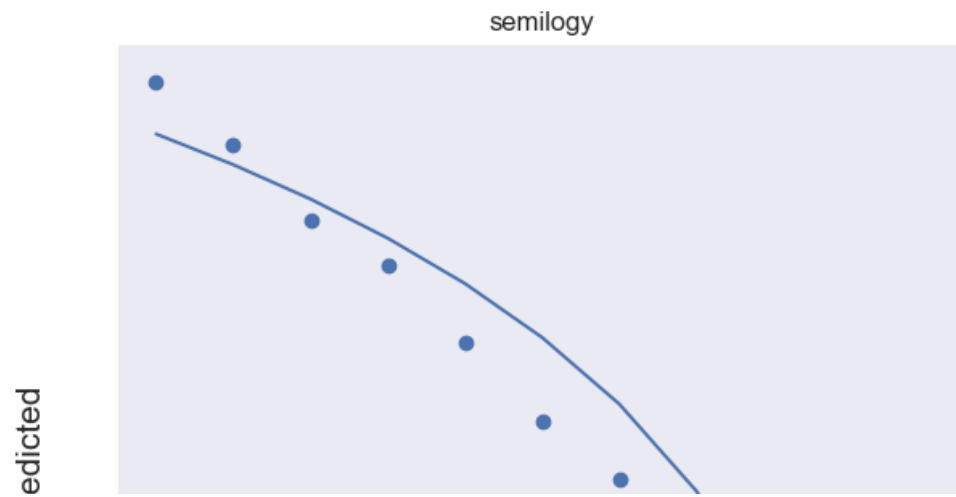


```
In [ ]: x_2 = df['x']
Y = [m*x_2+b, b*pow(x_2,m), b*np.exp(x_2*m), m*np.log(x_2) + b, 1/(m*x_2 + b)]
```



```
In [ ]: fig = plt.figure()
x_2 = df['x']
y = Y[0]
plt.subplot(2,2,1)
plt.scatter(df['x'], df['y'], label='Original Data')
plt.plot(x_2, y, label='y=mx+c')
plt.title('plot')
plt.xlabel('x', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
plt.grid()

plt.subplot(2,2,2)
plt.scatter(df['x'], df['y'], label='Original Data')
plt.semilogx(x_2, y, label='y=mx+c')
plt.title('semilogx')
plt.xlabel('x', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
plt.grid()

plt.subplot(2,2,3)
plt.scatter(df['x'], df['y'], label='Original Data')
plt.semilogy(x_2, y, label='y=mx+c')
plt.title('semilogy')
plt.xlabel('x', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
plt.grid()

plt.subplot(2,2,4)
plt.scatter(df['x'], df['y'], label='Original Data')
plt.loglog(x_2, y, label='y=mx+c')
plt.title('loglog')
plt.xlabel('x', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
plt.grid()
plt.show()
```





Plot

```
In [ ]: fig_1, axs = plt.subplots(2, 2)
# plt.rcParams['figure.figsize'] = [15, 15]
axs[0, 0].plot(x_2, Y[1])
axs[0, 0].scatter(df['x'], df['y'])
axs[0, 0].set_title('y=b*pow(x,m)')


axs[0, 1].plot(x_2, Y[2], 'tab:orange')
axs[0, 1].set_title('y=b*np.exp(x*m)')
axs[0,1].scatter(df['x'], df['y'])

axs[1, 0].plot(x_2, Y[3], 'tab:green')
axs[1, 0].set_title('y=m*np.log(x)')
axs[1,0].scatter(df['x'], df['y'])


axs[1, 1].plot(x_2, Y[4], 'tab:red')
axs[1, 1].set_title('y=1/(m*x + b)')
axs[1,1].scatter(df['x'], df['y'])

for ax in axs.flat:
    ax.set_xlabel('x', fontsize=15)
    ax.set_ylabel('y-predicted', fontsize=15)
    ax.grid()
```


$$y=b*\text{pow}(x,m)$$




$$y=b*\text{np.exp}(x*m)$$




$$y=m*\text{np.log}(x)$$



$$y=1/(m*x + b)$$





Semilogx





```
In [ ]: y = y_1
# plt.semilogx(x,y, marker='.', markersize=15, color='green')
fig_1, axs = plt.subplots(2, 2)
axs[0, 0].semilogx(x_2,Y[1], marker='.', markersize=15)
axs[0, 0].scatter(df['x'], df['y'])
axs[0, 0].set_title('y=b*pow(x,m)')



axs[0, 1].semilogx(x_2,Y[2], marker='.', markersize=15, color='green')
axs[0, 1].set_title('y=b*np.exp(x*m)')
axs[0,1].scatter(df['x'], df['y'])

axs[1, 0].semilogx(x_2,Y[3], marker='.', markersize=15, color='red')
axs[1, 0].set_title('y=m*np.log(x)')
axs[1,0].scatter(df['x'], df['y'])

axs[1, 1].semilogx(x_2,Y[4], marker='.', markersize=15, color='orange')
axs[1, 1].set_title('y=1/(m*x + b)')
axs[1,1].scatter(df['x'], df['y'])

for ax in axs.flat:
    ax.set_xlabel('x', fontsize=15)
    ax.set_ylabel('y-predicted', fontsize=15)
    ax.grid()
```



Semilogy


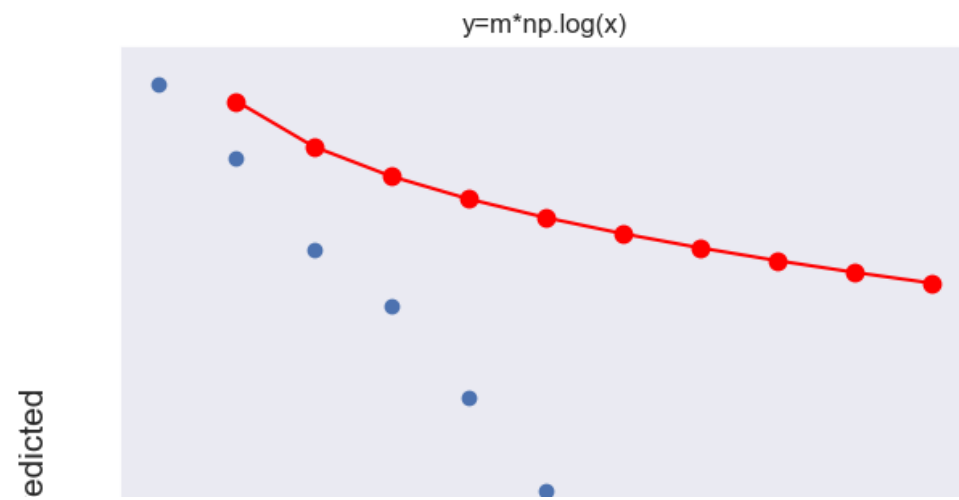
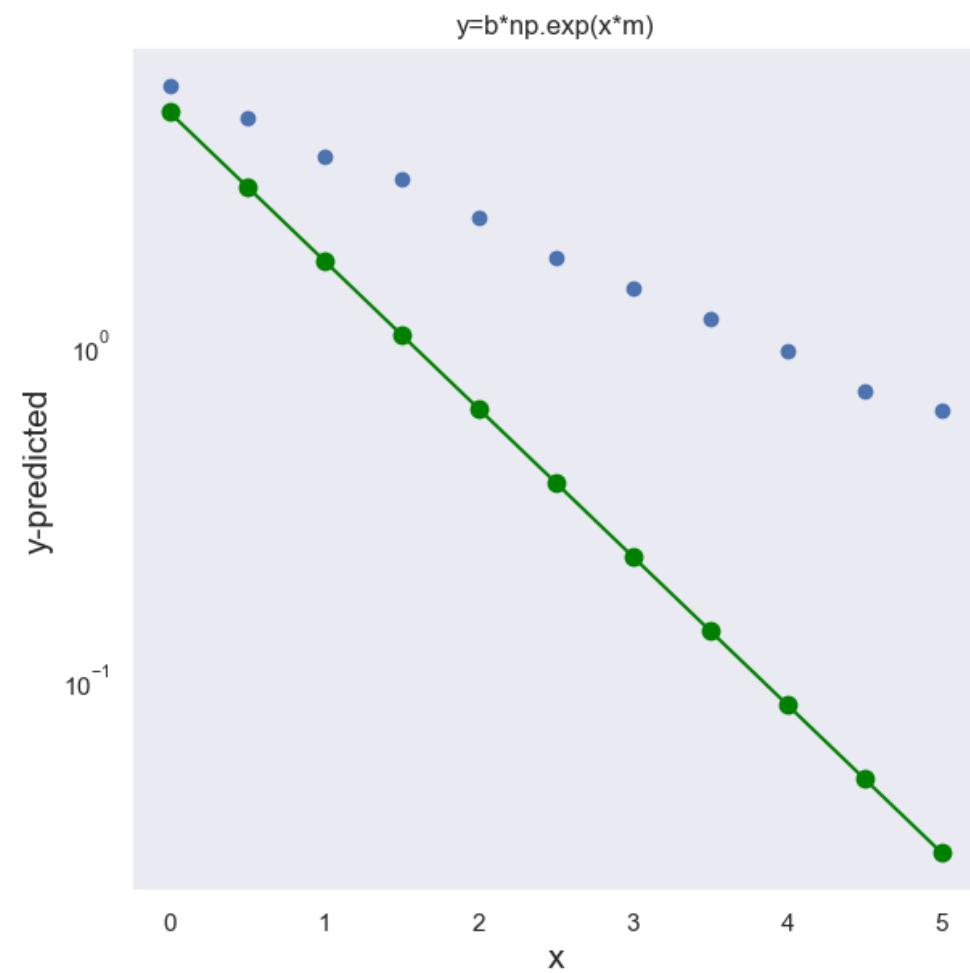

```
In [ ]: y = y_1
# plt.semilogx(x,y, marker='.', markersize=15, color='green')
fig_1, axs = plt.subplots(2, 2)
# plt.rcParams['figure.figsize'] = [15, 15]
axs[0, 0].semilogx(x_2, Y[1], marker='.', markersize=15)
axs[0, 0].scatter(df['x'], df['y'])
axs[0, 0].set_title('y=b*pow(x,m)')

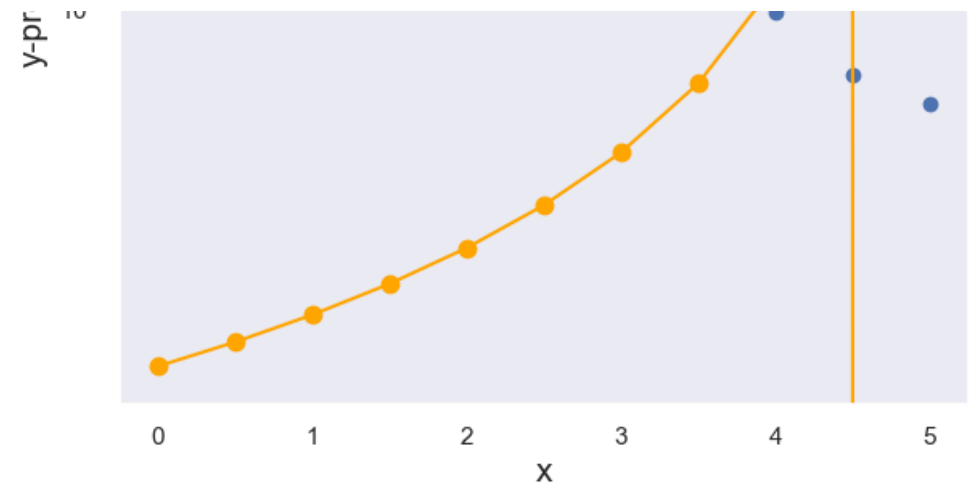

axs[0, 1].semilogx(x_2, Y[2], marker='.', markersize=15, color='green')
axs[0, 1].set_title('y=b*np.exp(x*m)')
axs[0, 1].scatter(df['x'], df['y'])

axs[1, 0].semilogx(x_2, Y[3], marker='.', markersize=15, color='red')
axs[1, 0].set_title('y=m*np.log(x)')
axs[1, 0].scatter(df['x'], df['y'])

axs[1, 1].semilogx(x_2, Y[4], marker='.', markersize=15, color='orange')
axs[1, 1].set_title('y=1/(m*x + b)')
axs[1, 1].scatter(df['x'], df['y'])

for ax in axs.flat:
    ax.set_xlabel('x', fontsize=15)
    ax.set_ylabel('y-predicted', fontsize=15)
    ax.grid()
```





loglog





```
In [ ]: y = y_1
# plt.semilogx(x,y, marker='.', markersize=15, color='green')
fig_1, axs = plt.subplots(2, 2)
axs[0, 0].loglog(x_2, Y[1], marker='.', markersize=15)
axs[0, 0].scatter(df['x'], df['y'])
axs[0, 0].set_title('y=b*pow(x,m)')



axs[0, 1].loglog(x_2, Y[2], marker='.', markersize=15, color='green')
axs[0, 1].set_title('y=b*np.exp(x*m)')
axs[0, 1].scatter(df['x'], df['y'])

axs[1, 0].loglog(x_2, Y[3], marker='.', markersize=15, color='red')
axs[1, 0].set_title('y=m*np.log(x)')
axs[1, 0].scatter(df['x'], df['y'])

axs[1, 1].loglog(x_2, Y[4], marker='.', markersize=15, color='orange')
axs[1, 1].set_title('y=1/(m*x + b)')
axs[1, 1].scatter(df['x'], df['y'])

for ax in axs.flat:
    ax.set_xlabel('x', fontsize=15)
    ax.set_ylabel('y-predicted', fontsize=15)
    ax.grid()
```





```
In [ ]: plt.rcParams['figure.figsize'] = [5, 5]
```

Question 3

```
In [ ]: df_3 = {'x': [0.5, 2.4, 3.2, 4.9, 6.5, 7.8], 'y': [0.8, 9.3, 37.9, 68.2, 155.0, 198.0]}
df_3 = pd.DataFrame(data=df_3)
x_3 = df_3['x'].to_numpy()
y_3 = df_3['y'].to_numpy()
```

```
In [ ]: plt.rcParams['figure.figsize'] = [8, 5]
```

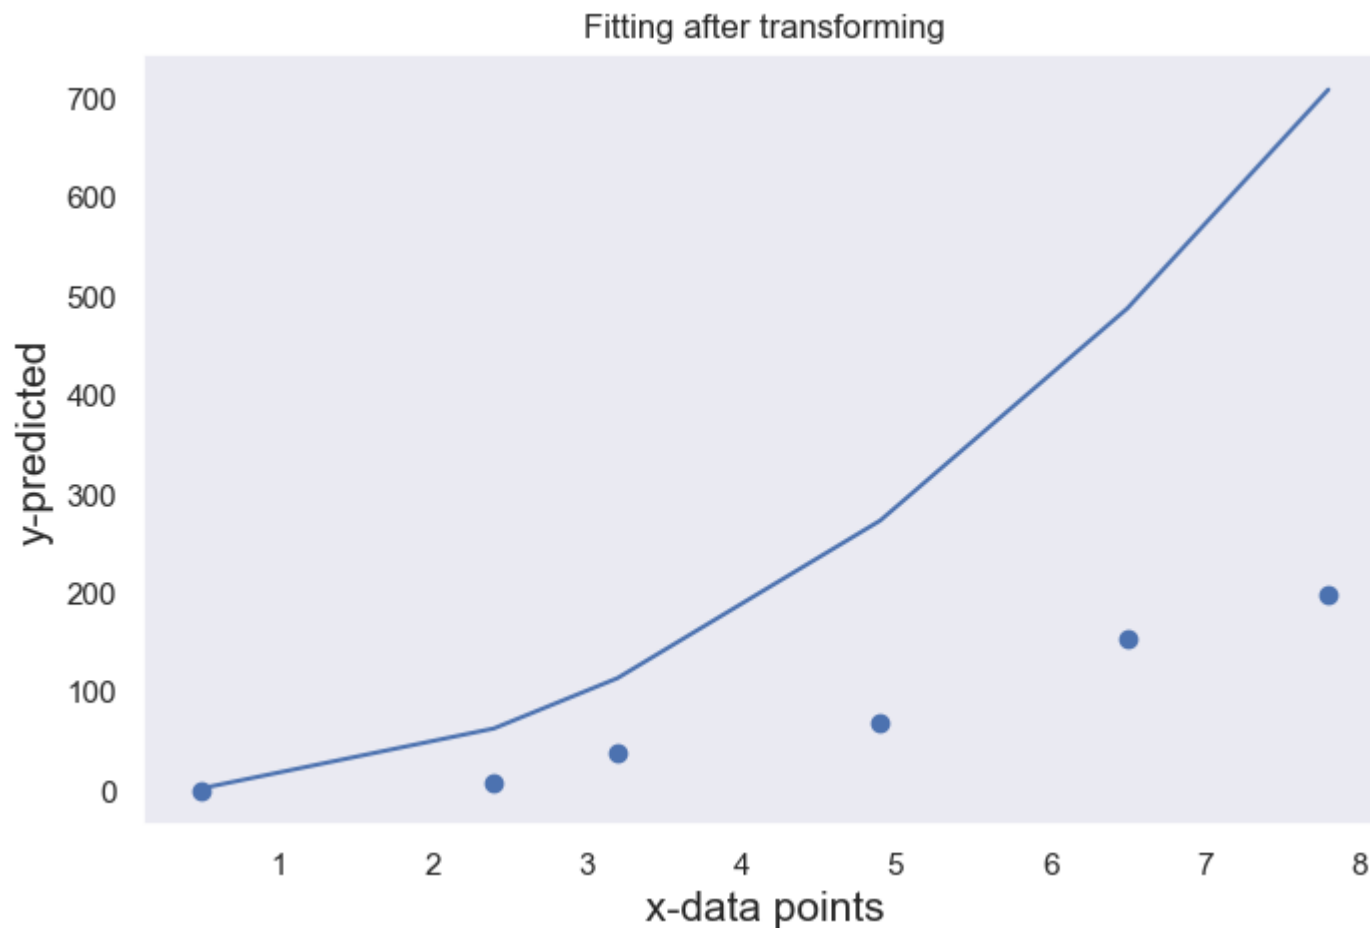
```
In [ ]: x_dash = np.log(x_3)
y_dash = np.log(y_3)
```

```
In [ ]: def powerfit(x, y):
    xm = np.mean(x)
    ym = np.mean(y)
    b1 = np.sum((x-xm)*(y-ym))/(np.sum((x-xm)**2))
    b0 = ym - b1*xm
    return (b1, b0)
b1, b0 = powerfit(x_dash, y_dash)
m = b1 #gradient
b = b0 #intercept
print('m: {}, b:{}'.format(b1, b0))
```

m: 2.049553636875363, b:1.022453728075393

```
In [ ]: b = pow(10,b)
```

```
In [ ]: fig = plt.figure()
y_predicted = b*pow(x_3,m)
# print(y_3)
plt.scatter(df_3['x'], df_3['y'], label='Original Data')
plt.plot(x_3, y_predicted, label='y=b*x^m')
# plt.title('Graph of y=-1.022x+5.005 and Original Data Points', fontsize=15)
plt.xlabel('x-data points', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
# plt.legend(loc="upper right", prop={'size':13})
plt.title('Fitting after transforming')
plt.grid()
plt.show()
```



NOT TRANSFORMING

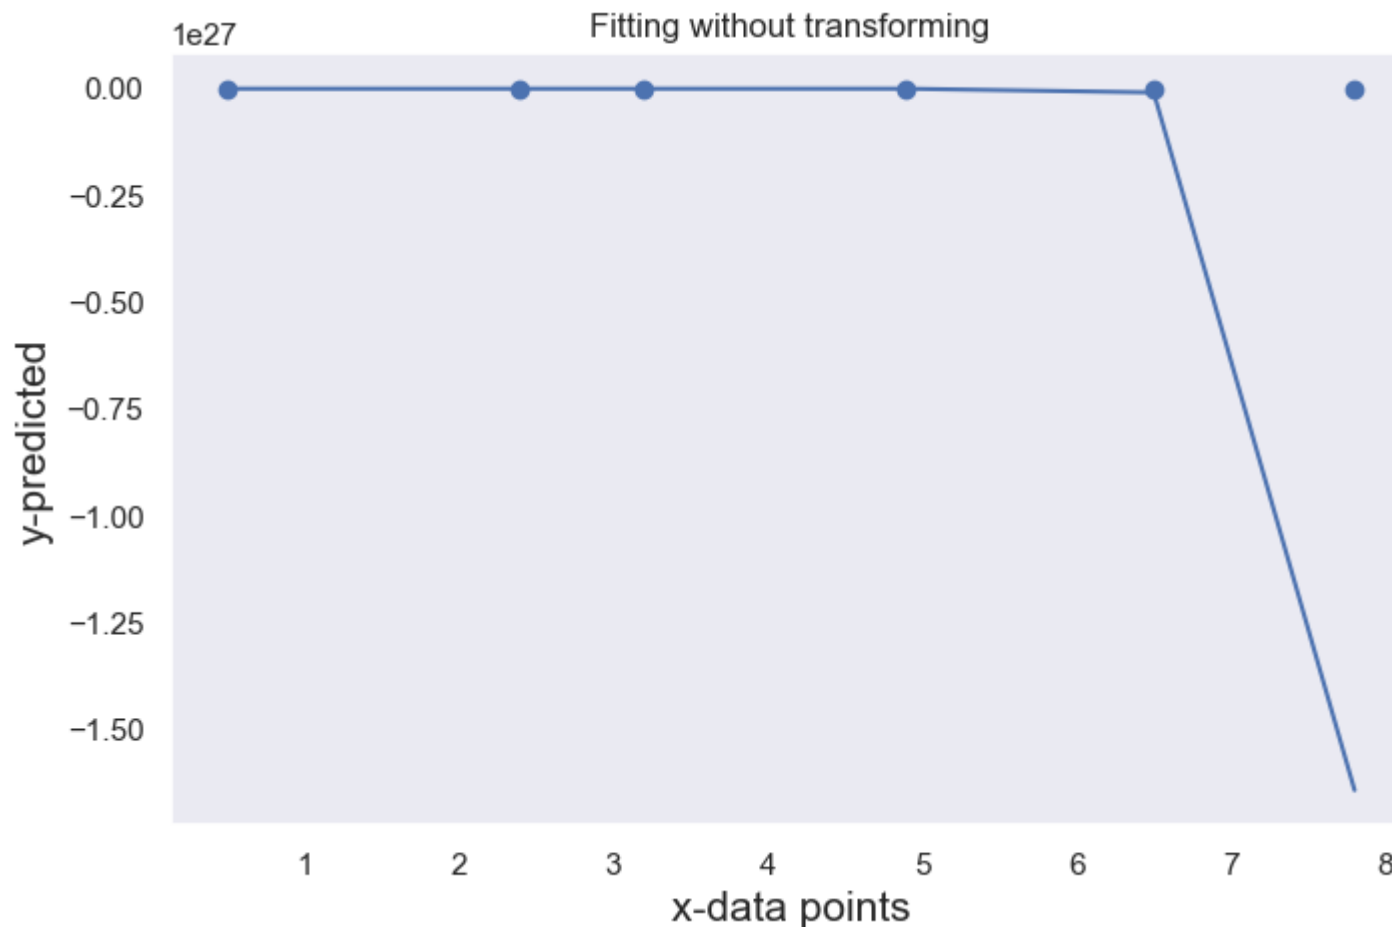
```
In [ ]: def powerfit(x, y):  
    xm = np.mean(x)  
    ym = np.mean(y)  
    b1 = np.sum((x-xm)*(y-ym))/(np.sum((x-xm)**2))  
    b0 = ym - b1*xm  
    return (b1, b0)  
b1, b0 = powerfit(x_3, y_3)  
m = b1 #gradient  
b = b0 #intercept  
print('m: {}, b:{}'.format(b1, b0))
```

m: 28.679241852643063, b:-42.73080314531158


```

In [ ]: fig = plt.figure()
y_predicted = b*pow(x_3,m)
# print(y_3)
plt.scatter(df_3['x'], df_3['y'], label='Original Data')
plt.plot(x_3, y_predicted, label='y=b*x^m')
# plt.title('Graph of y=-1.022x+5.005 and Original Data Points', fontsize=15)
plt.xlabel('x-data points', fontsize=15)
plt.ylabel('y-predicted', fontsize=15)
# plt.legend(loc="upper right", prop={'size':13})
plt.title('Fitting without transforming')
plt.grid()
plt.show()

```



Question 4

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np

fig = plt.figure()
ax = fig.add_subplot(projection='3d')

### Create the mesh in polar coordinates and compute corresponding Z.
r = np.linspace(0, 2)
theta = np.linspace(0, 2*np.pi)

R, THETA = np.meshgrid(r, theta)
Z = 4*R
# Express the mesh in the cartesian system.
X = R*np.cos(THETA)
Y = R*np.sin(THETA)

# Plot the surface.
ax.plot_surface(X, Y, Z, cmap=plt.cm.YlGnBu_r)

### top of cone
theta_t = np.linspace(0, 2*np.pi)
phi_t = np.linspace(0, np.pi/2)
# r_t = np.linspace(0, 2)
PHI_t, THETA_t = np.meshgrid(phi_t, theta_t)
# PHI_t, _ = np.meshgrid(phi_t, theta_t)

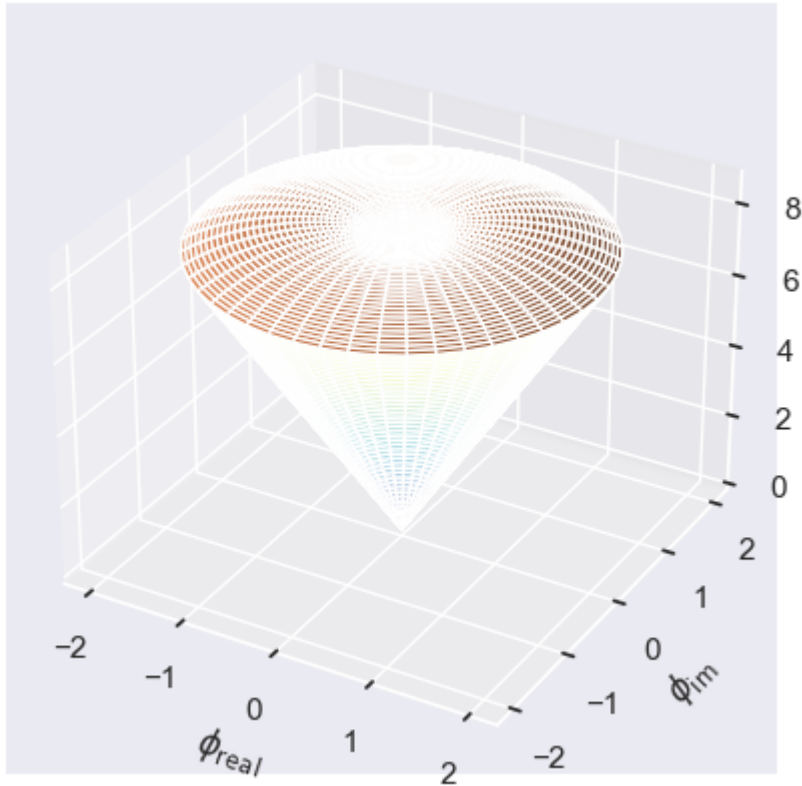
# Make data.
X_t = R*np.cos(THETA_t)*np.sin(PHI_t)
Y_t = R*np.sin(THETA_t)*np.sin(PHI_t)
Z_t = 8+R*np.cos(PHI_t)

# Plot the surface.
surf = ax.plot_surface(X_t, Y_t, Z_t)

# ax.view_init(25, 20)

# Tweak the limits and add latex math labels.
# ax.set_zlim(0, 1)
ax.set_xlabel(r' $\phi_{\mathrm{real}}$ ')
ax.set_ylabel(r' $\phi_{\mathrm{im}}$ ')
ax.set_zlabel(r' $V(\phi)$ ')
```

```
plt.show()
```



Exploratory Analysis on Real-World Data

Question 2 - Part 1

Load the dataset using pandas and display all necessary information contained in the file

```
In [ ]: pd.set_option('display.max_rows', None)
```

```
In [ ]: df = pd.read_csv("/Users/farjad.ahmed/Documents/Studies/ML Lab/Exercise_01/task1.txt")
```

```
In [ ]: df.columns
```

```
Out[ ]: Index(['Year', 'Month', 'DayofMonth', 'DayOfWeek', 'DepTime', 'CRSDepTime',  
              'ArrTime', 'CRSArrTime', 'UniqueCarrier', 'FlightNum', 'TailNum',  
              'ActualElapsedTime', 'CRSElapsedTime', 'AirTime', 'ArrDelay',  
              'DepDelay', 'Origin', 'Dest', 'Distance', 'TaxiIn', 'TaxiOut',  
              'Cancelled', 'CancellationCode', 'Diverted', 'CarrierDelay',  
              'WeatherDelay', 'NASDelay', 'SecurityDelay', 'LateAircraftDelay'],  
             dtype='object')
```

```
In [ ]: dframe.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99260 entries, 0 to 99259
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  99260 non-null  int64
1   Month                 99260 non-null  int64
2   DayofMonth            99260 non-null  int64
3   DayOfWeek             99260 non-null  int64
4   DepTime               97847 non-null  float64
5   CRSDepTime            99260 non-null  int64
6   ArrTime               97693 non-null  float64
7   CRSArrTime            99260 non-null  int64
8   UniqueCarrier         99260 non-null  object
9   FlightNum             99260 non-null  int64
10  TailNum               98156 non-null  object
11  ActualElapsedTime     97659 non-null  float64
12  CRSElapsedTime        99249 non-null  float64
13  AirTime               97659 non-null  float64
14  ArrDelay              97659 non-null  float64
15  DepDelay              97847 non-null  float64
16  Origin                99260 non-null  object
17  Dest                  99260 non-null  object
18  Distance              99260 non-null  int64
19  TaxiIn                97693 non-null  float64
20  TaxiOut               97841 non-null  float64
21  Cancelled             99260 non-null  int64
22  CancellationCode      1420 non-null   object
23  Diverted              99260 non-null  int64
24  CarrierDelay          19747 non-null  float64
25  WeatherDelay          19747 non-null  float64
26  NASDelay              19747 non-null  float64
27  SecurityDelay         19747 non-null  float64
28  LateAircraftDelay     19747 non-null  float64
dtypes: float64(14), int64(10), object(5)
memory usage: 22.0+ MB

```

```
In [ ]: dframe.head(10)
```

Out []:	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	UniqueCarrier	FlightNum	...	TaxiIn	TaxiOut	Cancelled	Canc
0	2008	1	1	2	120.0	1935	309.0	2130	9E	5746	...	3.0	18.0	0	
1	2008	1	1	2	555.0	600	826.0	835	AA	1614	...	7.0	11.0	0	
2	2008	1	1	2	600.0	600	728.0	729	YV	2883	...	7.0	16.0	0	
3	2008	1	1	2	601.0	605	727.0	750	9E	5743	...	4.0	12.0	0	
4	2008	1	1	2	601.0	600	654.0	700	AA	1157	...	5.0	10.0	0	
5	2008	1	1	2	636.0	645	934.0	932	NW	1674	...	11.0	22.0	0	
6	2008	1	1	2	646.0	655	735.0	750	CO	340	...	6.0	15.0	0	
7	2008	1	1	2	650.0	700	841.0	857	XE	541	...	6.0	11.0	0	
8	2008	1	1	2	650.0	650	1139.0	1145	AA	1182	...	4.0	12.0	0	
9	2008	1	1	2	654.0	700	1117.0	1133	B6	1060	...	13.0	13.0	0	

10 rows × 29 columns

Question 2 - part 2

You are tasked as a data scientist to create a story that is visually appealing from this data. Create plots using matplotlib/seaborn that will depict such interesting stories from flights that depart from and arrive in the Austin region. The figures should be annotated properly and also easily understandable on the first glance. A list of questions that can be explored/answered as reference are given below. Of course, you are free to explore any other possibilities.

Investigate what time of the day it is best to fly so as to have the least possible delays. Does this change with airlines?

```
In [ ]: #Adding a column that holds total delays in a journey except the arrival delay, this is added to the dframe1 as 'total_delays'
dframe1 = dframe
# dframe1['total_delays'] = dframe1['DepDelay'] + dframe1['CarrierDelay'] + dframe1['WeatherDelay'] + dframe1['NASDelay'] + dfr
# # dframe1 = dframe1[dframe1.total_delays.notnull()]
# dframe1 = dframe1.dropna(subset=['CRSDepTime'])
```

```
In [ ]: dframe1['avg_delays'] = dframe1[['DepDelay', 'CarrierDelay', 'WeatherDelay', 'NASDelay', 'SecurityDelay', 'LateAircraftDelay']]
# dframe1.head(50)
```

```
In [ ]: #Converting single digit hour values in CRSDepTime to hhmm by padding a 0 at the beginning
dframe1['CRSDepTime'] = dframe1['CRSDepTime'].astype(int).astype(str)
```

```
dframe1['CRSDepTime'] = dframe1['CRSDepTime'].str.zfill(4)
#Splitting the column to obtain hh (hours) values from the dataframe which will be used to group by the data.
#Minutes are ignored, hence the values will be accurate to hourly time periods
dframe1['DT_Hours'] = dframe1['CRSDepTime'].str.slice(0,2)
```


```
In [ ]: result1 = dframe1.groupby('DT_Hours', as_index=False)['avg_delays'].min()
result1
```

```
Out[ ]:
```

	DT_Hours	avg_delays
0	00	-10.0
1	05	-15.0
2	06	-17.0
3	07	-36.0
4	08	-29.0
5	09	-17.0
6	10	-23.0
7	11	-22.0
8	12	-18.0
9	13	-20.0
10	14	-19.0
11	15	-42.0
12	16	-22.0
13	17	-20.0
14	18	-18.0
15	19	-22.0
16	20	-23.0
17	21	-20.0
18	22	-14.0
19	23	-14.0

```
In [ ]: ax = sns.lineplot(result1, x=result1['DT_Hours'], y=result1['avg_delays'])
```

```
ax.set(xlabel='Time of the Day in Hours', ylabel='Minimum Flying Delay', title='Delay vs Hours')
plt.show()
```



Investigating whether the flying delays vary with airlines

```
In [ ]: result2 = dframe1.groupby(['DT_Hours', 'UniqueCarrier'], as_index=False)['avg_delays'].min()
# result2
```


```
In [ ]: # Tried doing this with sql for exploring purposes, hence some are done with pandas functions some with sql
# query = "Select DT_Hours, UniqueCarrier, min(avg_delays) as min_delays from dframe1 group by DT_Hours, UniqueCarrier"
# dframe2 = sqldf(query, globals())
# dframe2.head(50)
```

```
In [ ]: fig, ax = plt.subplots(figsize=(20, 10))
```



```
sns.scatterplot(dframe1.groupby([dframe1['UniqueCarrier'], 'DT_Hours'])['avg_delays'].mean().unstack(), linewidth=0.5, ax=ax)

plt.xlabel('UniqueCarrier')
plt.ylabel('Average Delay')
plt.xticks(rotation=450)
plt.legend()
plt.show()
```



```
In [ ]: q = "Select DT_Hours, UniqueCarrier, min(avg_delays) as least_delays from result2 group by DT_Hours"
df_day = sqldf(q, globals())
df_day
```


```
Out[ ]:
```

	DT_Hours	UniqueCarrier	least_delays
0	00	B6	-10.0
1	05	YV	-15.0
2	06	YV	-17.0
3	07	YV	-36.0
4	08	YV	-29.0
5	09	CO	-17.0
6	10	MQ	-23.0
7	11	YV	-22.0
8	12	B6	-18.0
9	13	B6	-20.0
10	14	XE	-19.0
11	15	YV	-42.0
12	16	AA	-22.0
13	17	AA	-20.0
14	18	XE	-18.0
15	19	AA	-22.0
16	20	MQ	-23.0
17	21	YV	-20.0
18	22	CO	-14.0
19	23	AA	-14.0

Better graphical representation of best time of flying with minimum delays along with the carrier

```
In [ ]: #Bar chart representation of the best time to fly on a given day, w.r.t the flight carrier
ax = sns.barplot(data=df_day, x=df_day['DT_Hours'], y=df_day['least_delays'], ci = None)
for container, number in zip(ax.containers, df_day.UniqueCarrier):
    ax.bar_label(container, labels=list(df_day['UniqueCarrier']))
ax.set(title='Best flying time and the Flight Carrier')
```

Out[]: [Text(0.5, 1.0, 'Best flying time and the Flight Carrier')]



Investigate what time of the year it is more suited to fly so as to have the delays minimum and does the destination affect this? You can lay insights on some popular destinations for the task.


Affects of time of the year

```
In [ ]: # q = "Select Month, Dest, min(avg_delays) as least_delays from dframe1 group by Month"
# result2 = sqldf(q, globals())
# result2.head(12)
result2 = dframe1.groupby('Month')['avg_delays'].min().reset_index(name='least_delays')
result2
```

Out[]:


	Month	least_delays
0	1	-29.0
1	2	-22.0
2	3	-22.0
3	4	-22.0
4	5	-36.0
5	6	-20.0
6	7	-42.0
7	8	-20.0
8	9	-20.0
9	10	-17.0
10	11	-26.0
11	12	-16.0

```
In [ ]: sns.set(rc={'figure.figsize':(32,10)})
ax1 = sns.lineplot(data=result2, x=result2['Month'], y=result2['least_delays'])
ax1.set(xlabel='Time of the Year in Months', ylabel='Minimum Flying Delay', title='Delay vs Month')
plt.show()
```



```
In [ ]: fig, ax = plt.subplots(figsize=(30, 10))
# fig = plt.figure()
sns.scatterplot(dframe1.groupby([dframe1['UniqueCarrier'], 'Month'])['avg_delays'].mean().unstack(), linewidth=0.5, ax=ax)

plt.xlabel('UniqueCarrier')
plt.ylabel('Average Delay')
plt.xticks(rotation=450)
plt.legend()
plt.show()
```



```
In [ ]: q = "Select Month, Dest, min(avg_delays) as Least_Delay from dframe1 group by Month, Dest"
mydf = sqldf(q, globals())
# mydf.head(10)
# ax = sns.barplot(data=mydf, x=mydf['Month'], y=mydf['Least_Delay'], ci = None)
# for container, number in zip(ax.containers, mydf.Dest):
#     ax.bar_label(container, labels = set(list(mydf.Dest)))
# ax.set(title='Time of the Year to Fly with Minimum Delay w.r.t Flight Carriers')
# plt.show()

#Bar chart representation of the best time to fly on a given day, w.r.t the flight carrier
ax = sns.barplot(data=mydf, x=mydf['Month'], y=mydf['Least_Delay'], ci = None)
for container, number in zip(ax.containers, mydf.Dest):
    ax.bar_label(container, labels=list(set(list(mydf['Dest']))))
ax.set(title='Best flying time and the Flight Carrier')
```

AttributeError Traceback (most recent call last)


Cell In [191], line 13

```
11 ax = sns.barplot(data=mydf, x=mydf['Month'], y=mydf['Least_Delay'], ci = None)
12 for container, number in zip(ax.containers, mydf.Dest):
--> 13     ax.bar_label(container, labels=list(set(list(mydf['Dest']))))
14 ax.set(title='Best flying time and the Flight Carrier')
```

File ~/Library/Python/3.9/lib/python/site-packages/matplotlib/axes/_axes.py:2712, in Axes.bar_label(self, container, labels, fmt, label_type, padding, **kwargs)

```
2707 annotations = []
2709 for bar, err, dat, lbl in itertools.zip_longest(
2710     bars, errs, datavalues, labels
2711 ):
-> 2712     (x0, y0), (x1, y1) = bar.get_bbox().get_points()
2713     xc, yc = (x0 + x1) / 2, (y0 + y1) / 2
2715     if orientation == "vertical":
```

AttributeError: 'NoneType' object has no attribute 'get_bbox'



You can lay insights on some popular destinations for the task.

In []: *#Finding the most popular destination*

```
dframe1.groupby('Dest')['Dest'].count().sort_values(ascending=False).head(1)
```

```
Out [ ]: Dest
AUS      49637
Name: Dest, dtype: int64
```

Explore some airports that are bad to fly to. Does the time of day or year affect this?

```
In [ ]: dframe1.head(5)
```

```
Out [ ]:
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	UniqueCarrier	FlightNum	...	Cancelled	CancellationCode	Div
0	2008	1	1	2	120.0	1935	309.0	2130	9E	5746	...	0		NaN
1	2008	1	1	2	555.0	0600	826.0	835	AA	1614	...	0		NaN
2	2008	1	1	2	600.0	0600	728.0	729	YV	2883	...	0		NaN
3	2008	1	1	2	601.0	0605	727.0	750	9E	5743	...	0		NaN
4	2008	1	1	2	601.0	0600	654.0	700	AA	1157	...	0		NaN

5 rows × 31 columns

```
In [ ]: dframe1['avg_delayWithArrivalDelay'] = dframe1[['DepDelay', 'CarrierDelay', 'WeatherDelay', 'NASDelay', 'SecurityDelay', 'LateA  
dframe1.head(10)
```


Out[]:	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	UniqueCarrier	FlightNum	...	CancellationCode	Diverted	Carr
0	2008	1	1	2	120.0	1935	309.0	2130	9E	5746	...	NaN	0	
1	2008	1	1	2	555.0	0600	826.0	835	AA	1614	...	NaN	0	
2	2008	1	1	2	600.0	0600	728.0	729	YV	2883	...	NaN	0	
3	2008	1	1	2	601.0	0605	727.0	750	9E	5743	...	NaN	0	
4	2008	1	1	2	601.0	0600	654.0	700	AA	1157	...	NaN	0	
5	2008	1	1	2	636.0	0645	934.0	932	NW	1674	...	NaN	0	
6	2008	1	1	2	646.0	0655	735.0	750	CO	340	...	NaN	0	
7	2008	1	1	2	650.0	0700	841.0	857	XE	541	...	NaN	0	
8	2008	1	1	2	650.0	0650	1139.0	1145	AA	1182	...	NaN	0	
9	2008	1	1	2	654.0	0700	1117.0	1133	B6	1060	...	NaN	0	


10 rows × 32 columns

```
In [ ]: df_day = dframe1.groupby(['Dest'])['avg_delayWithArrivalDelay'].max().reset_index(name='max_delays')
df_day.head(10)
```

Out[]:	Dest	max_delays
0	ABQ	118.000000
1	ATL	395.857143
2	AUS	220.714286
3	BNA	74.714286
4	BOS	187.428571
5	BWI	115.285714
6	CLE	157.571429
7	CLT	94.285714
8	CVG	157.142857
9	DAL	251.000000

```
In [ ]: sns.set(rc={'figure.figsize':(35,10)})
```

```
ax = sns.lineplot(x=df_day['Dest'], y=df_day['max_delays'], ci=None)
ax.set(xlabel='Destination', ylabel='Maximum Delay', title='Airports vs Max Delays')
# ax.set_xticks()
plt.show()
```



```
In [ ]: # #Bar chart representation of the best time to fly on a given day, w.r.t the flight carrier
# sns.set(rc={'figure.figsize':(50,20)})
# ax = sns.barplot(data=df_day, x=df_day['Dest'], y=df_day['Least_Delay'], ci = None)
# for container, number in zip(ax.containers, df_day.Dest):
#     ax.bar_label(container)
# ax.set(title='Best flying time and the Flight Carrier')
# plt.show()
```

```
In [ ]: dframe1.head(10)
```

Out[]:

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	UniqueCarrier	FlightNum	...	CancellationCode	Diverted	Carr
0	2008	1	1	2	120.0	1935	309.0	2130	9E	5746	...	NaN	0	
1	2008	1	1	2	555.0	0600	826.0	835	AA	1614	...	NaN	0	
2	2008	1	1	2	600.0	0600	728.0	729	YV	2883	...	NaN	0	
3	2008	1	1	2	601.0	0605	727.0	750	9E	5743	...	NaN	0	
4	2008	1	1	2	601.0	0600	654.0	700	AA	1157	...	NaN	0	
5	2008	1	1	2	636.0	0645	934.0	932	NW	1674	...	NaN	0	
6	2008	1	1	2	646.0	0655	735.0	750	CO	340	...	NaN	0	
7	2008	1	1	2	650.0	0700	841.0	857	XE	541	...	NaN	0	
8	2008	1	1	2	650.0	0650	1139.0	1145	AA	1182	...	NaN	0	
9	2008	1	1	2	654.0	0700	1117.0	1133	B6	1060	...	NaN	0	

10 rows × 32 columns

In []:

```
#Time of the day affects
df_day = dframe1.groupby(['Dest', 'DT_Hours'], as_index=False)['avg_delayWithArrivalDelay'].max().sort_values(by='avg_delayWithArrivalDelay')
df_day.rename(columns = {'avg_delayWithArrivalDelay':'max_delays'}, inplace = True)
df_day.head(10)
```

Out[]:

	Dest	DT_Hours	max_delays
4	ATL	05	395.857143
208	LAS	19	286.000000
144	FLL	14	280.428571
77	DAL	09	251.000000
289	ORD	17	228.000000
23	AUS	12	220.714286
29	AUS	18	213.000000
27	AUS	16	209.428571
253	MDW	16	209.000000
31	AUS	20	202.142857

```
In [ ]: d = list(set(list(df_day.Dest)))
```

```
In [ ]: #Bar chart representation of the best time to fly on a given day, w.r.t the flight carrier
var = list(set(list(df_day.Dest)))
sns.set(rc={'figure.figsize':(50,20)})
ax = sns.barplot(data=df_day, x=df_day['DT_Hours'], y=df_day['max_delays'], ci = None)
for container, number in zip(ax.containers, df_day.Dest):
    ax.bar_label(container, labels=d)
ax.set(title='Worst Airports to Fly w.r.t Months')
plt.show()
```

```
-----
AttributeError                                Traceback (most recent call last)
```


```
Cell In [1124], line 6
```

```
    4 ax = sns.barplot(data=df_day, x=df_day['DT_Hours'], y=df_day['max_delays'], ci = None)
    5 for container, number in zip(ax.containers, df_day.Dest):
----> 6     ax.bar_label(container, labels=d)
    7 ax.set(title='Worst Airports to Fly w.r.t Months')
    8 plt.show()
```

```
File ~/Library/Python/3.9/lib/python/site-packages/matplotlib/axes/_axes.py:2712, in Axes.bar_label(self, container, labels, format, label_type, padding, **kwargs)
```

```
    2707 annotations = []
    2709 for bar, err, dat, lbl in itertools.zip_longest(
    2710     bars, errs, datavalues, labels
    2711 ):
-> 2712     (x0, y0), (x1, y1) = bar.get_bbox().get_points()
    2713     xc, yc = (x0 + x1) / 2, (y0 + y1) / 2
    2715     if orientation == "vertical":
```

```
AttributeError: 'NoneType' object has no attribute 'get_bbox'
```



```
In [ ]: #Time of the year affects
df_month = dframe1.groupby(['Dest', 'Month'], as_index=False)['avg_delayWithArrivalDelay'].max().sort_values(by='avg_delayWithA
df_month.rename(columns = {'avg_delayWithArrivalDelay':'max_delays'}, inplace = True)
df_month.head(10)
```

```
Out[ ]:
```

	Dest	Month	max_delays
20	ATL	12	395.857143
244	LAS	2	286.000000
162	FLL	7	280.428571
100	DAL	3	251.000000
380	ORD	8	228.000000
27	AUS	7	220.714286
30	AUS	10	213.000000
28	AUS	8	209.428571
319	MDW	2	209.000000
32	AUS	12	202.142857

```
In [ ]: d = list(set(list(df_month.Dest)))
```

```
In [ ]: #Bar chart representation of the best time to fly on a given day, w.r.t the flight carrier
var = list(set(list(df_month.Dest)))
sns.set(rc={'figure.figsize':(50,20)})
ax = sns.barplot(data=df_day, x=df_month['Month'], y=df_month['max_delays'], ci = None)
for container, number in zip(ax.containers, df_month.Dest):
    ax.bar_label(container, labels=d)
ax.set(title='Worst Airports to Fly w.r.t Months')
plt.show()
```

AttributeError Traceback (most recent call last)


Cell In [1127], line 6

```
 4 ax = sns.barplot(data=df_day, x=df_month['Month'], y=df_month['max_delays'], ci = None)
 5 for container, number in zip(ax.containers, df_month.Dest):
----> 6     ax.bar_label(container, labels=d)
 7 ax.set(title='Worst Airports to Fly w.r.t Months')
 8 plt.show()
```

File ~/Library/Python/3.9/lib/python/site-packages/matplotlib/axes/_axes.py:2712, in Axes.bar_label(self, container, labels, fmt, label_type, padding, **kwargs)

```
2707 annotations = []
2709 for bar, err, dat, lbl in itertools.zip_longest(
2710     bars, errs, datavalues, labels
2711 ):
-> 2712     (x0, y0), (x1, y1) = bar.get_bbox().get_points()
2713     xc, yc = (x0 + x1) / 2, (y0 + y1) / 2
2715     if orientation == "vertical":
```

AttributeError: 'NoneType' object has no attribute 'get_bbox'



Investigate on how the pattern of flights to various destinations alter over the course of year.


```
In [ ]: #Pattern of flights over the course of a year along with Flight
result1 = dframe1.groupby(['Month', 'Dest', 'FlightNum'])['FlightNum'].count().reset_index(name='CountOfFlights')
result1.head(10)
```


Out[]:

	Month	Dest	FlightNum	CountOfFlights
0	1	ABQ	311	31
1	1	ABQ	315	27
2	1	ATL	466	2
3	1	ATL	470	2
4	1	ATL	1254	1
5	1	ATL	1535	11
6	1	ATL	1590	30
7	1	ATL	3906	2
8	1	ATL	4325	25
9	1	ATL	4338	31

	Month	Dest	FlightNum	CountOfFlights
0	1	ABQ	311	31
1	1	ABQ	315	27
2	1	ATL	466	2
3	1	ATL	470	2
4	1	ATL	1254	1
5	1	ATL	1535	11
6	1	ATL	1590	30
7	1	ATL	3906	2
8	1	ATL	4325	25
9	1	ATL	4338	31

```
In [ ]: fig, ax = plt.subplots(figsize=(20, 10))
Months = ['January', 'February', 'March', 'April', 'May', 'June', 'July', 'August', 'September', 'October', 'November', 'December']
sns.scatterplot(result1.groupby(
[result1['Dest'], 'Month'])['CountOfFlights'].mean().unstack(), linewidth=0.5, ax=ax)
# sns.barplot(TYResult.groupby(
# [TYResult['Dest'], 'Month'])['Min_Delay'].mean().unstack(), linewidth=0.5, ax=ax)
plt.xlabel('Destination')
plt.ylabel('Count Of Flights')
plt.title('Patterns of Flights Over the Course of Year')
plt.xticks(rotation=45)
plt.legend(Months)
plt.show()
```



Question 2

In this part we will examine the data containing information on every Olympic medallist that is listed by participant count in top 20 sports, dating back to 1896. Load the dataset task2.txt and perform statistical analysis on the dataset. Specifically, do the following:

```
In [ ]: df = pd.read_csv("/Users/farjad.ahmed/Documents/Studies/ML Lab/Exercise_01/task2.txt")
```

```
In [ ]: #Displaying all the necessary information from the file
df.head()
```

Out []:

	id	name	sex	age	height	weight	team	noc	games	year	season	city	sport	event	medal
0	16	Juhamatti Tapio Aaltonen	M	28	184	85.0	Finland	FIN	2014 Winter	2014	Winter	Sochi	Ice Hockey	Ice Hockey Men's Ice Hockey	Bronze
1	17	Paavo Johannes Aaltonen	M	28	175	64.0	Finland	FIN	1948 Summer	1948	Summer	London	Gymnastics	Gymnastics Men's Individual All-Around	Bronze
2	17	Paavo Johannes Aaltonen	M	28	175	64.0	Finland	FIN	1948 Summer	1948	Summer	London	Gymnastics	Gymnastics Men's Team All-Around	Gold
3	17	Paavo Johannes Aaltonen	M	28	175	64.0	Finland	FIN	1948 Summer	1948	Summer	London	Gymnastics	Gymnastics Men's Horse Vault	Gold
4	17	Paavo Johannes Aaltonen	M	28	175	64.0	Finland	FIN	1948 Summer	1948	Summer	London	Gymnastics	Gymnastics Men's Pommel Horse	Gold

Compute the 95th percentile of heights for the competitors in all Athletic events for gender Female. Note that sport refers to the broad sports (Athletics) and event is the specific event (100-meter sprint).

```
In [ ]: df_1 = df[(df['sex']=='F') & (df['sport']=='Athletics') ]
df_1['height'].quantile(0.95)
```

Out []: 183.0

Find the single woman’s event that depicts the highest variability in the height of the competitor across the entire history of Olympics. Use the standard deviation as the yardstick for this.

```
In [ ]: df_2 = df[(df['sex']=='F')]
#This could be done in two ways
#1
df_2.groupby('event')['height'].std().sort_values(ascending=False).reset_index(name='std').head(1)
```

Out []:

	event	std
0	Rowing Women's Coxed Fours	10.86549

```
In [ ]: #2
df_2 = df_2.groupby('event')['height'].std()
df_2 = df_2.to_frame()
df_2['height'].idxmax()
```

```
Out [ ]: "Rowing Women's Coxed Fours"
```

We wish to know how the average age of swimmers in Olympic has evolved with time. How has this changed over time? Does the trend for this differs from male to female? It will be easy to create a data frame that will allow one to visualise these trends with time. Plot a line graph that depicts separate line for male and female competitors. The plot must have a caption that is informative enough to answer the 2 questions that have been asked in this part.

```
In [ ]: #Creating two dataframes, one for females and the other for males
df_F = df[df['sex']=='F']
df_M = df[df['sex']=='M']
```

```
In [ ]: df_M1 = df_M.groupby('year')['age'].agg('mean')
df_M1 = df_M1.to_frame()
# df_M1
```

```
In [ ]: df_F1 = df_F.groupby('year')['age'].agg('mean')
df_F1 = df_F1.to_frame()
# df_F1
```

Here I see this is not the best method of doing this, I will have to join these two dataframes for female avg ages and male avg ages. Instead I found an easier method by using unstack to do this, shown below.

```
In [ ]: df_s = df[(df['sport']=='Swimming')]
```

```
In [ ]: # This is an easy way of doing this
sns.set(rc={'figure.figsize':(20,10)})
out = df_s.groupby(['year','sex'])['age'].agg('mean').unstack()
out.plot(xlabel='year', ylabel='Avg Age', title='Avg Age of Female and Male Swimmers wrt Time')
plt.show()
```

Avg Age of Female and Male Swimmers wrt Time

